

# K-MEANS CLUSTERING

-APPLIED MULTIVARIATE ANALYSIS & STATISTICAL  
LEARNING-

MMA Chapter 15.4.1, ISLR 10.3.1

Lecturer: Darren Homrighausen, PhD

# Preamble:

- State K-means
- Discuss how it can be computed
- Define a way of choosing the number of clusters

# CLUSTERING INTRODUCTION

We will focus on two particular clustering algorithms

- **K-MEANS:** Seeks to partition the the observations into a pre-specified number of clusters.
- **HIERARCHICAL:** Produces a tree-like representation of the observations, known as a **dendrogram**.

There are advantages (disadvantages) to both approaches.

We can cluster **observations** on the basis of the **features** in order to find subgroups of observations.

Just as easily, we can find clusters of **features** based on the **observations** to find subgroups in the features.

We will focus on clustering the observations. You can cluster features by transposing  $\mathbb{X}$  (that is, clustering on  $\mathbb{X}^T$ ).

# K-MEANS

1. Select a number of clusters  $K$ .
2. Let  $C_1, \dots, C_K$  partition  $\{1, 2, 3, \dots, n\}$  such that
  - ▶ All observations belong to some set  $C_j$ .
  - ▶ No observation belongs to more than one set.
3. K-means attempts to form these sets by making **within-cluster variation**,  $W(C_k)$ , as small as possible.

$$\min_{C_1, \dots, C_K} \sum_{k=1}^K W(C_k).$$

4. To Define  $W$ , we need a concept of distance. By far the most common is Euclidean

$$W(C_k) = \frac{1}{|C_k|} \sum_{i, i' \in C_k} \|X_i - X_{i'}\|_2^2.$$

That is, the average (Euclidean) distance between all cluster members.

# K-MEANS

It turns out

$$\min_{C_1, \dots, C_K} \sum_{k=1}^K W(C_k) \quad (1)$$

is too hard of a problem to solve computationally ( $K^n$  partitions!).

So, we make a greedy approximation:

1. Randomly assign observations to the  $K$  clusters
2. Iterate until the cluster assignments stop changing:
  - ▶ For each of  $K$  clusters, compute the **centroid**, which is the  $p$ -length vector of the means in that cluster.
  - ▶ Assign each observation to the cluster whose centroid is closest (in Euclidean distance).

This procedure is guaranteed to decrease (1) at each step.

**Warning:** It finds only a local minimum, not necessarily the global one. Which local min depends on step 1.

# K-MEANS: A SUMMARY

To fit K-means, you need to

1. Pick  $K$  (inherent in the method)
2. Convince yourself you have found a good solution (due to the randomized approach to the algorithm).

It turns out that 1. is difficult to do in a principled way. We will discuss this shortly

For 2., a commonly used approach is to run K-means many times with different starting points. Pick the solution that has the smallest value for

$$\min_{C_1, \dots, C_K} \sum_{k=1}^K W(C_k)$$

As an aside, why can't we use the approach for picking  $K$ ?

## K-MEANS: A SUMMARY

To fit K-means, you need to

1. Pick  $K$  (inherent in the method)
2. Convince yourself you have found a good solution (due to the randomized approach to the algorithm).

It turns out that 1. is difficult to do in a principled way. We will discuss this shortly

For 2., a commonly used approach is to run K-means many times with different starting points. Pick the solution that has the smallest value for

$$\min_{C_1, \dots, C_K} \sum_{k=1}^K W(C_k)$$


As an aside, why can't we use the approach for picking  $K$ ?

(We would choose  $K = n$ )

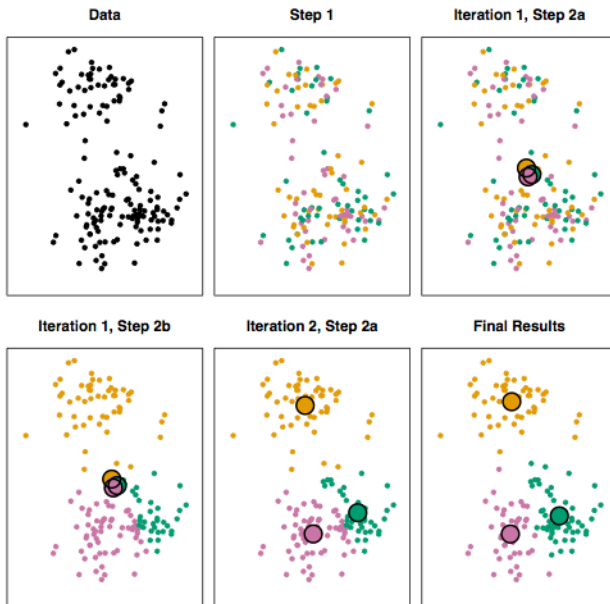
# K-means



# K-MEANS: VARIOUS $K$ 'S



# K-MEANS: ALGORITHM AT WORK



# K-MEANS: FINDING GOOD LOCAL MINIMUM



# K-MEANS IN R

Like usual, the interface with R is very basic

```
n    = 30
X1 = rnorm(n)
X2 = rnorm(n)
X    = cbind(X1,X2)
K = 3
kmeans.out = kmeans(X, centers=K)
> names(kmeans.out)
[1] "cluster"      "centers"      "totss"      "withinss"
[5] "tot.withinss" "betweenss"    "size"
> kmeans.out$cluster
[1] 2 2 2 2 2 2 1 1 2 2 3 1 2 1 2 2 2
3 1 2 2 1 3 2 1 3 3 1 2 3
```


# K-MEANS IN R

Like usual, the interface with R is very basic

```
n      = 30
X1 = rnorm(n)
X2 = rnorm(n)
X     = cbind(X1,X2)
K = 3
kmeans.out = kmeans(X, centers=K)
> names(kmeans.out)
[1] "cluster"      "centers"      "totss"      "withinss"
[5] "tot.withinss" "betweenss"    "size"
> kmeans.out$cluster
 [1] 2 2 2 2 2 2 1 1 2 2 3 1 2 1 2 2 2
 3 1 2 2 1 3 2 1 3 3 1 2 3

plot(X, col=(kmeans.out$cluster+1), xlab="", ylab="",
      pch=20, cex=2)
```

# K-MEANS IN R



# K-MEANS IN R

## Another example

```
x = matrix(rnorm(50*2),ncol=2)
x[1:25,1] = x[1:25,1] + 3
x[1:25,2] = x[1:25,2] -4

kmeans.out = kmeans(x,centers=2,nstart=20)
```

# K-MEANS IN R




FIGURE: Two clusters (which is the true number)



# K-MEANS IN R




FIGURE: Three clusters

## K-MEANS IN R: COMPARISON

R provides several objects in the `kmeans` output.

$W(C_k)$  is the same as: `kmeans(x,centers=K)$withinss`

$\sum_{k=1}^K W(C_k)$  is the same as: `kmeans(x,centers=K)$tot.withinss`

```
> kmeans(x,centers=4,nstart=1)$tot.withinss
[1] 19.12722
> kmeans(x,centers=4,nstart=20)$tot.withinss
[1] 18.5526
> kmeans(x,centers=5,nstart=20)$tot.withinss
[1] 12.01893
```

# Choosing $K$

# CHOOSING THE NUMBER OF CLUSTERS

Why is it important?

- It might make a big difference (concluding there are  $K = 2$  cancer sub-types versus  $K = 3$ )
- One of the major goals of statistical learning is automatic inference. A good way of choosing  $K$  is certainly a part of this

# CHOOSING THE NUMBER OF CLUSTERS

Sometimes, the number of clusters is fixed ahead of time:

- Segmenting a client database into  $K$  clusters for  $K$  salesmen
- Compressing an image using vector quantization ( $K$  is the compression rate)

Most of the time, it isn't so straight forward. Why is this a hard problem?

- Determining the number of clusters is **hard** (for humans) unless the data is low dimensional
- It is just as hard to explain what we are looking for  
(ie: in classification, we want a classifier that predicts well. In clustering, we want a clusterer to ... what?)

## REMINDER: WHAT DOES $K$ -MEANS DO?

Given a number of clusters  $K$ , we (approximately) minimize:

$$\sum_{k=1}^K W(C_k) = \sum_{k=1}^K \frac{1}{|C_k|} \sum_{i, i' \in C_k} \|x_i - x_{i'}\|_2^2.$$


We can rewrite this in terms of the **centroids** as

$$W(K) = 2 \sum_{k=1}^K \sum_{i \in C_k} \|x_i - \bar{x}_k\|_2^2,$$

## MINIMIZING $W$ IN $K$

Of course, a lower value of  $W$  is better. Why not minimize  $W$ ?

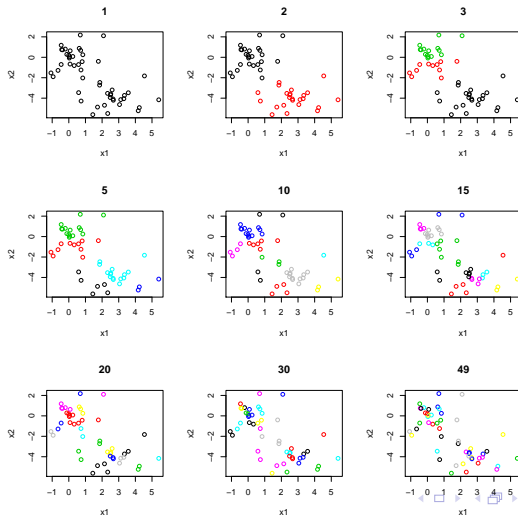
```
plotW = rep(0,49)
for(K in 1:49){
  plotW[K] = kmeans(x,centers=K,nstart=20)$tot.withinss
}
```



# MINIMIZING $W$ IN $K$

Of course, a lower value of  $W$  is better. Why not minimize  $W$ ?

A look at the cluster solution





## BETWEEN-CLUSTER VARIATION

Within-cluster variation measures how **tightly grouped** the clusters are. As we increase  $K$ , this will always decrease


What we are missing is **between-cluster variation**, ie: how spread apart the groups are

$$B = \sum_{k=1}^K |C_k| \|\bar{X}_k - \bar{X}\|_2^2,$$

where  $|C_k|$  is the number of points in  $C_k$ , and  $\bar{X}$  is the grand mean of all observations:

$$\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$$

# BETWEEN AND WITHIN CLUSTER VARIATION EXAMPLE




$$B = |C_1| \|\bar{X}_1 - \bar{X}\|_2^2 + |C_2| \|\bar{X}_2 - \bar{X}\|_2^2$$

$$W = \sum_{i \in C_1} |C_1| \|\bar{X}_1 - X_i\|_2^2 + \sum_{i \in C_2} |C_2| \|\bar{X}_2 - X_i\|_2^2$$

## CAN WE JUST MAXIMIZE $B$ ?

Sadly, no. Just like  $W$  can be made arbitrarily small,  $B$  will always be increasing with increasing  $K$ .



## CH INDEX

Ideally, we would like our cluster assignment to **simultaneously** have small  $W$  and large  $B$ .

This is the idea behind **CH index**. For clustering assignments coming from  $K$  clusters, we record  $CH$  score:

$$CH(K) = \frac{B(K)/(K-1)}{W(K)/(n-K)}$$

To choose  $K$ , pick some maximum number of clusters to be considered ( $K_{\max} = 20$ , for example) and choose the value of  $K$  that

$$\hat{K} = \arg \max_{K \in \{2, \dots, K_{\max}\}} CH(K).$$

**Note:**  $CH$  is undefined for  $K = 1$ .


```
ch.index = function(x,kmax,iter.max=100,nstart=10,  
                    algorithm="Lloyd")  
{  
  ch = numeric(length=kmax-1)  
  n = nrow(x)  
  for (k in 2:kmax) {  
    a = kmeans(x,k,iter.max=iter.max,nstart=nstart,  
              algorithm=algorithm)  
    w = a$tot.withinss  
    b = a$betweenss  
    ch[k-1] = (b/(k-1))/(w/(n-k))  
  }  
  return(list(k=2:kmax,ch=ch))  
}
```

# A SIMULATED EXAMPLE


```
x = matrix(rnorm(50*2),ncol=2)
x[1:25,1] = x[1:25,1] + 3
x[1:25,2] = x[1:25,2] -4
```

We want to cluster this data set using K-means with  $K$  chosen via  $CH$  index.

# CH PLOT



# CORRESPONDING SOLUTION





# Postamble:

- State K-means  
(K-means partitions the feature space, creating clusters)
- Discuss how it can be computed  
(K-means can be computed by randomly setting centroids and iteratively re-assigning points to clusters and recomputing centroids)
- Define a way of choosing the number of clusters  
(The CH index is a way to choose the number of clusters)